

Monte Carlo Simulation and Analysis of Gacha Systems

Alex Hoang

CECS 381 Stochastic Computing Final Project, Spring 2026

May 10, 2026

1 Problem Description and Motivation

Gacha games are a genre of video games where players can spend in-game currency on randomized "pulls" to obtain characters and equipment. Two notable examples are Wuthering Waves and Genshin Impact. Typically, there are two types of banners: a featured character banner and a featured weapon banner. Each banner has its own "pity" system, which increases the chances of getting a 5-star item as you pull more times without success. This project builds a model of a gacha banner system loosely based on Wuthering Waves and Genshin Impact, and uses Monte Carlo simulation to analyze two questions:

1. How does the simulation compare with an analytical analysis of the gacha system?
2. How does shifting the "soft-pity" onset change the distribution of pulls a player needs to obtain a full kit from a featured banner set (a character and a weapon banner)?

2 Model Setup

A gacha system typically has two independent banners: a character banner and a weapon banner. Pity counters don't transfer between them, so pulling on one banner doesn't affect the other. They are independent. Table 1 will list the variables that are relevant.

Table 1: Relevant Variables

Value Type	Symbol	Value
Base 5-star rate (both banners)	p_0	0.008
Hard pity threshold (both)	n_h	80 pulls
Soft-pity onset, character (swept)	n_s^c	{50, 60, 70}
Soft-pity onset, weapon (fixed)	n_s^w	65
50/50 win probability, character	q	0.5
Featured rate, weapon		1.0

Soft pity ramp. This mechanic applies to each banner independently. When used, we will plug in that banner's respective n_s from Table 1. We will use n to be the number of pulls since the last 5-star. The per-pull 5-star chance is p_0 for the first $n_s - 1$ pulls. Once we hit the soft-pity onset, the probability begins to ramp up linearly until it hits 1 (which is 100% probability) at pull n_h .

We can write this as a piecewise function:

$$p_n = \begin{cases} p_0 & 1 \leq n < n_s \\ p_0 + \frac{n - n_s}{n_h - n_s}(1 - p_0) & n_s \leq n \leq n_h. \end{cases} \quad (1)$$

In other words, when we hit n_h , we will always have a 5-star pull with no exceptions.

Character-banner 50/50 system. It should be noted that getting a 5-star on the character banner doesn't mean that you will get the featured character. There's a 50/50 chance it's actually some other standard 5-star character instead. Most gacha games include a guarantee system to ensure that players can get what they want even if they are unlucky. If you lose the 50/50, then the next 5-star on that banner is guaranteed to be the featured one. We can actually look at this like a two-state Markov chain between {Fifty-Fifty, Guarantee}.

- In **Fifty-Fifty**: the player has a 50-50 chance of getting the featured character or a standard 5-star character in this state. Thus, the probability $q = 0.5$. If the player rolls the featured character, they stay in Fifty-Fifty. If not, they move to the Guarantee state.
- In **Guarantee**: the next 5-star will be the featured character upon a successful 5-star roll. Upon acquisition, the player moves back to Fifty-Fifty.

It should be noted that this system only applies to the character banner. The weapon banner uses the same soft-pity ramp from equation (1) with its own onset n_s^w . However, it has no fifty-fifty system: every 5-star on the weapon banner is the featured weapon.

Full-kit problem. The player will start with no 5-star acquisitions, whether it is the character or the weapon. The player will pull on the character banner first until they get the featured character. Then they will switch to the weapon banner until they get the featured weapon. We will track the number of pulls using variable T .

3 Simulation Methodology

Each pull uses one $\text{Uniform}(0, 1)$ draw against the soft-pity ramp from equation (1) to decide whether it's a 5-star. If we are in the Fifty-Fifty state, there will be another draw to determine if we get a standard or featured character. Effectively, the second draw is a Bernoulli trial. We will run $N = 10^5$ trials of the full kit problem. The number of trials was chosen as to give us a large enough sample size to get a good estimation for the expected values.

Analytical reference. For a given banner, let N be the number of pulls until the first 5-star. The probability that this happens on exactly pull n is the chance of failing the first $n - 1$ pulls and then succeeding on the n -th. Since the pulls are independent, this is just a product:

$$P(N = n) = (1 - p_1)(1 - p_2) \dots (1 - p_{n-1})(p_n).$$

An easier way to look at this would be to use product notation. From there, we can plug the new equation into $\mathbb{E}[X] = \sum x P(X = x)$ where X is N :

$$P(N = n) = p_n \prod_{i=1}^{n-1} (1 - p_i), \quad \mathbb{E}[N] = \sum_{n=1}^{n_h} n P(N = n). \quad (2)$$

We actually can and will compute this using python loops to get an analytical estimation. From there, we can use it as a reference to check the simulator. For now, we will use N_c to represent the result of applying these equations to the character banner.

Due to the Fifty-Fifty mechanic, we will use Y to represent the number of pulls to get the featured character specifically. There are also something else to consider: the player's starting state. If we start in the Guarantee state, the next 5-star is always featured, so $\mathbb{E}[Y \mid \text{Guarantee}] = \mathbb{E}[N_c]$. If we start in Fifty-Fifty, we have to consider the total expectation of rolling.

This means that there are two scenarios:

- With probability q , we may win the 50/50 on our first 5-star which takes $\mathbb{E}[N_c]$ pulls on average.
- With probability $1 - q$ (which should be the same), we may lose the 50/50. This moves us to the Guarantee state, and need a second, independent cycle to get the guaranteed 5-star. This means that the expected number of pulls should be $2 \mathbb{E}[N_c]$. Do note that the first and second cycles are independent.

By the law of total expectation, conditioning on whether we win or lose the first 50/50, the overall expected number of pulls is the weighted sum of the two conditional expectations:

$$\mathbb{E}[Y \mid \text{Fifty-Fifty}] = P(\text{win}) \cdot \mathbb{E}[Y \mid \text{win}] + P(\text{lose}) \cdot \mathbb{E}[Y \mid \text{lose}].$$

Substituting $P(\text{win}) = q$, $P(\text{lose}) = 1 - q$, and the conditional expectations from the two bullets above:

$$\mathbb{E}[Y \mid \text{Fifty-Fifty}] = q \mathbb{E}[N_c] + (1 - q)(2 \mathbb{E}[N_c]) = (2 - q) \mathbb{E}[N_c] = \frac{3}{2} \mathbb{E}[N_c] \quad \text{when } q = \frac{1}{2}. \quad (3)$$

This is the formula the simulator should reproduce in Q1.

Confidence intervals. For our analysis we need to find the confidence intervals. To do so, we will use the central limit theorem:

$$\bar{Y} \pm 1.96 \cdot \frac{\hat{\sigma}}{\sqrt{n}}, \quad (4)$$

where $\hat{\sigma}$ is the sample standard deviation across the $n = 10^5$ trials and 1.96 is the standard normal critical value for the two-tailed 95% level. This will be calculated in the simulation.

4 Research Questions and Results

4.1 Q1: How does the simulation compare to an analytical analysis?

We ran 10^5 trials of the character banner from each starting state, Fifty-Fifty and Guarantee, and compared the sample means against the formulas in equation (3). 95% CIs are computed using equation (4).

Table 2: Simulator vs. analytical conditional means at baseline $n_s^c = 60$.

Starting state	Simulated $\bar{Y}$	95% CI	Analytical Expected Value Y
Fifty-Fifty	76.628	[76.403, 76.852]	76.498
Guarantee	50.972	[50.842, 51.103]	50.999




Figure 1: The running mean of $\hat{\mathbb{E}}[Y \mid \text{Fifty-Fifty}]$ converging as more trials are added. The horizontal axis is on log scale, and the dashed line shows the analytical value.

Result. Both simulated means (Fifty-Fifty and Guarantee) land inside the 95% confidence interval around the analytical value, which is a good indicator that the simulator is correct and that we can trust our values.

4.2 Q2: How does soft-pity onset change the full-kit distribution?

Keeping everything else the same, we will try three different character-banner soft-pity onsets: $n_s^c \in \{50, 60, 70\}$. Earlier onset means the ramp starts sooner, so the 5-star probability hits high values earlier in the pull sequence. For each onset we still simulate 10^5 trials.

Table 3: Total pulls T to complete the kit, by soft-pity onset. 10^5 trials each; 95% CIs in brackets, computed using equation (4).

n_s^c	Mean	95% CI	$Q_{0.50}$	$Q_{0.90}$	$Q_{0.99}$
50	121.99	[121.75, 122.23]	124	179	190
60	129.86	[129.60, 130.13]	133	195	206
70	137.31	[137.02, 137.60]	141	211	221




Figure 2: Distribution of total full-kit pulls under each soft-pity onset condition. Earlier onset shifts the whole distribution left.

Result. From what we find, the earlier the onset, the more the distribution will shift left. It would also appear that the tail tends to shrink a little bit more than the mean. Moving n_s^c from 70 to 50 drops the mean by about 15 pulls, or roughly 11%. We also find that the 99th percentile drops significantly from 221 to 190 if we move the onset from 70 to 50. This is a large drop of about 14%, which is 31 pulls. Considering that the tail shrinks a bit faster than the mean, reducing the onset might be an effective way to help balance the experience for unlucky players without affecting the average player too much.

5 Limitations and Future Work

In our simulation, we use a linear ramp-up for the soft-pity system. Although this simulation was supposed to simulate a gacha system, we should note that not all gacha games use a linear soft-pity system. So the values and analysis we have found may not necessarily be true in all real production gacha games. A future work could analyze different soft-pity curves and see how that changes the pull distribution.

It should also be noted that most players will not start at exactly 0 initial pulls on a banner. Players are actually able to save up pulls on one feature banner and the pity from the pull count will transfer over to the next featured banner when the current one ends and the new one replaces it. I mention this because this analysis paper does not entirely capture how the player would really approach getting a featured character and their weapon in-game. There is actually some strategy involved when planning to pull on a banner depending on what the player may want. A future work could potentially numerically analyze an optimal strategy for a player to maximize how much they can get out of a particular featured banner.